

Date of publication xxxx 00, 0000, date of current version xxxx 00, 0000.

Digital Object Identifier 10.1109/ACCESS.2017.DOI

Fraud Detection on Medicare Claims Using SQL-Based Feature Engineering and Machine Learning

ZAYED M. UDDIN¹, MD. ABU BAKAR SIDDIK MISHKAT², ULUGBEK ABDURAKHMONOV³, SARIKA RAJEEV^{4*}, AND TIMOTHY OLADUNNI^{5*}

¹Department of Mathematics and Computer Science, Coppin State University, Baltimore, MD 21216, USA (e-mail: zayed.md.uddin273@gmail.com)

²Department of Mathematics and Computer Science, Coppin State University, Baltimore, MD 21216, USA (e-mail: mishkatfamily@gmail.com)

³Department of Mathematics and Computer Science, Coppin State University, Baltimore, MD 21216, USA (e-mail: ulugbekabdurakhmonov.tennis@gmail.com)

⁴Department of Mathematics and Computer Science, Coppin State University, Baltimore, MD 21216, USA (e-mail: srjeev@coppin.edu)

⁵Department of Computer Science, Morgan State University, Baltimore, MD 21251, USA (e-mail: Timothy.Oladunni@morgan.edu)

Corresponding authors: Dr. Sarika Rajeev* (e-mail: srjeev@coppin.edu) and Dr. Timothy Oladunni* (e-mail: timothy.oladunni@morgan.edu).

ABSTRACT Healthcare fraud in Medicare always results in substantial financial losses and is becoming harder to detect as claims data become increasingly sophisticated and imbalanced. In this paper, we present a complete pipeline combining an SQL-based feature store, supervised binary classification, and unsupervised clustering to detect and categorize fraudulent activity. Unsupervised clustering is applied only to ground-level fraudulent cases. Then we classified the fraudulent providers and clustered the high-risk providers into separate fraud type groups. Using inpatient, outpatient, and beneficiary tables from a Medicare dataset, we engineered reusable provider- and claim-level features for model training and evaluation. This included rule-based indicators for duplicate claims, upcoding, and overcharging, as well as temporal aggregates and peer-normalized z-scores. We evaluated multiple supervised learning models to classify fraudulent providers; a regularized logistic regression model achieved the strongest performance with an AUPRC of 0.73 and an ROC-AUC of 0.94. Furthermore, we applied the HDBSCAN algorithm, Principal Component Analysis (PCA), and association rule mining to known fraudulent cases and high-risk providers. This unsupervised approach allowed us to cluster and characterize distinct fraud typologies successfully. In general, this integrated methodology establishes a robust evaluation benchmark and demonstrates an effective, scalable approach to profiling complex fraud schemes in healthcare.

INDEX TERMS Association rule mining, HDBSCAN fraud clustering, Healthcare insurance fraud types, Medicare fraud detection, SQL-feature engineering, Tree-based ensemble learning, Unsupervised anomaly detection.

I. INTRODUCTION

HEALTHCARE fraud and abuse divert limited resources from patient care and impose substantial financial losses on a large public insurance program such as Medicare. It is estimated that 3-10% of total healthcare expenses are subject to fraud and waste each year, amounting to tens to hundreds of billions of dollars globally and, on average, 300 billion USD annually in the US [1]-[3]. This waste also eroded trust in health systems [3], [4]. Reviews show that as health insurance costs rise in many countries, fraudulent activities are also increasing proportionally, compromising patient safety and raising questions about the quality of care [2],

[4]. Medicare claims provide a lot of information regarding patients, providers, diagnoses, procedures, deductibles, and reimbursements for detection. However, these features are highly imbalanced and multidimensional [1], [2], [5]. With the multidimensional dataset, detection is sophisticated, as fraud cases are rare among all claims [1]-[3], [6]-[10]. Rule-based systems are easy to implement but offer less scope for automation in fraud detection and can lead to high false-positive rates [2]-[4]. These limitations have inspired comprehensive research in healthcare to invent effective strategies to prevent fraud [1]-[4], [11], [12].

The primary objective of this study is to develop and

evaluate a comprehensive, hybrid pipeline for detecting and characterizing Medicare fraud. By integrating an SQL-based feature store with supervised classical machine learning models and unsupervised clustering algorithms (HDBSCAN), this research aims to establish a scalable benchmark to overcome the limitations of traditional rule-based systems and highly imbalanced claims data.

To achieve our primary objective, this research introduces the following contributions to the field of healthcare fraud detection:

- **Scalable SQL-Based Feature Engineering:** We develop a domain-specific feature store that translates raw claims data into peer-normalized indicators such as temporal aggregates, upcoding and overcharging metrics, effectively resolving the class imbalance of Medicare data.
- **Classical Machine Learning Benchmark:** We show that when equipped with robust feature engineering, models like logistic regression can match or exceed the predictive power of tree-based ensembles, achieving a ROC-AUC of 0.94 and AUPRC of 0.73.
- **Unsupervised Fraud Typology Mapping:** Moving beyond binary classification, we introduce a methodology combining HDBSCAN clustering and association rule mining to successfully isolate and label real-world fraud schemes such as high-revenue inpatient overcharging (Cluster 1) versus high-volume outpatient overbilling (Cluster 0 and 2).

II. RELATED WORK

Recent studies present a variety of audit-based checks and ML techniques for fraud detection in health insurance, including statistical models, ensemble learning, support vector machines, deep learning, and hybrid pipelines [1]-[4]. Du Preez et al. did not identify a single best approach in their review of machine learning methods [1]. The main reasons are differences in dataset features, different types of fraud, and evaluation techniques [1]. Having access to the ground truth is also a challenge here [1]. Kumaraswamy et al. worked on data-mining methods for fraud detection [2]. They focused primarily on the significant financial loss, evolution of fraud schemes, and differences in evaluation between academic models and real-world systems [2]. Thaifur et al. reported the discovery of various types of fraudulent activities and focused only on providers [4]; they skipped inspections at the claim level. Razzaq and Shah highlighted challenges related to imbalanced data, scarcity of trustworthy labels, privacy concerns, and the practicality of the models in terms of usability [3].

Supervised models were used to classify fraudulent claims and providers. Nabrawi and Alanazi evaluated several supervised models on Saudi insurance claims, and their work showed that tree-based models and the SVM model performed well when the ground truth for fraud cases was available [13]. Hamid et al. studied the CMS DE-SynPUF claims and found that gradient boosting performed well despite class imbalance and that provider-level features were carefully implemented

[11]. Liu and Vasarhelyi proposed a clustering model using geographic and payment data [14]. Their approach combined claim and provider features to flag suspicious cases [14]. Kumar and Sountharajan built an optimized deep learning model with proper hyperparameter tuning for fraud detection and noted challenges with feature engineering and class imbalance [15]. Similarly, F. Ahammed et al. developed a deep neural network framework that integrated feature selection (Chi-squared method) and data sampling techniques (SMOTE) to address the challenges of high dimensionality and class imbalance in Medicare claims data [16]. Yet their model struggles to meet the computational resources required for real-time detection on live-streaming datasets and relies on provider-level feature engineering [10]. Hancock et al. proposed explainable ensemble models for Medicare data using gradient-boosted trees; however, their reliance on tree depth and feature reduction lacks the transparency of explicit rule-based indicators [17].

Unsupervised anomaly detection methods are found to be more suitable when working with limited labels for fraud cases. Bairy et al. tested Isolation Forest, Local Outlier Factor, KNN, and autoencoders on a healthcare dataset; their methods can spot rare anomalies that need further review [5]. De Meulemeester used deep embeddings to flag anomalies and added necessary explanations for the investigators [18]. Kanksha et al. flagged fraud cases without having enough fraud labels [19]. Settipalli and Gangadharan worked on WMTDBC, which converted correlated claim data into independent variables [20]. They also used distance-based detection, which improved the detection of multivariate extremities [20]. Massi et al. first grouped providers into peer clusters and then identified each outlier within each cluster [12]. Saldamli et al. worked on blockchain systems to enhance transaction security and decentralize claims data [21]. Matloob et al. proposed a two-module framework utilizing sequence analysis, a rule engine, and a prediction-based engine to identify healthcare insurance fraud [22]. Their initial data validation was hindered by the difficulty of acquiring confidential employee insurance data [22].

While supervised models achieve high predictive performance [11], [13], [15], they struggle in the absence of ground-truth labels and often lack claim-level granularity [4]. Furthermore, benchmark ensemble methods, which are highly accurate, often sacrifice the clinical transparency required for auditing (Hancock et al. [17]). Conversely, unsupervised methods [5], [12], [19], [20] successfully bypass the need for labels but flag broad anomalies that overwhelm investigators and lack predictive probability scoring. With scalable SQL-based feature engineering, supervised provider classification, and unsupervised clustering, our approach mitigates class imbalance and produces interpretable, actionable fraud typologies. All of these works are briefly summarized in Table 1.

III. METHODOLOGY

The complete sequence of operations for our hybrid approach from SQL-based data curation to unsupervised fraud typology

TABLE 1. Summary of State-of-the-Art Methodologies and Research Gaps

References	Methodology	Strengths	Research Gaps
Thaifur et al. [4]	Fraud Discovery	Successfully categorizes different types of fraudulent activities	Focuses exclusively on the provider level, completely skipping critical inspections at the individual claim level
Bairy et al. [5] & Kanksha et al. [19]	Unsupervised Models: Isolation Forest and Local Outlier Factor (LOF) algorithms	Best performance in Isolation Forest & Handling Complex Data	Need for Advanced Architectures, Limited to a Single Dataset
F. Ahammed et al. [10]	Chi-squared Method & SMOTE	Reduces model complexity and overfitting, highly effective at minimizing false negatives	Model struggles with the computational resources required for real-time detection on live-streaming datasets
Hamid et al. [11]	Association Rule Mining & Unsupervised Models	Fewer False Positives, Adaptable & Robust System	Lack of representation of the diversity and complexity of real-world fraud scenarios
Settipalli & Gangadharan [20] & Massi et al. [12]	WMTDBC & Peer Clustering	Successfully isolates outliers within peer groups (e.g., multivariate extremities)	Distance and peer-based models identify extremes but lack a supervised scoring mechanism to assign definitive fraud probabilities
Nabrawi & Alanazi [13]	Supervised Models (Tree-based & SVM)	Achieves strong classification performance on Saudi insurance claims	Small sample size and a limited number of features, which restricted its ability to cover all aspects of fraud
Liu & Vasarhelyi [14]	Geographic & Payment Clustering	Spots concurrent provider fraud, subscriber fraud & collusive fraud	Flags “suspicious” claims only for further investigation and cannot definitively prove a claim is fraudulent on its own
Kumar & Sountharajan [15]	Novel Deep Learning Framework: EHOA-CNN-12	Balanced Architecture, Improved Optimization & High Accuracy	High Computational Requirements, Lack of Interpretability, Lack of Adversarial Testing
Hancock et al. [17]	Ensemble Supervised Feature Selection Technique	Massive Dimensionality Reduction, Computational Efficiency	Lack of Domain Diversity, Reliance on Retrospective Labels
De Meulemeester [18]	Categorical Embeddings	Utilizes embeddings to discover unknown fraud & highly interpretable	Conducted purely for research purposes on historical data & has not yet been implemented for active, operational goals
Saldamli et al. [21]	Blockchain Systems	Provides advanced security and decentralization for transactions	Focuses on data security infrastructure rather than the algorithmic detection of sophisticated billing anomalies
Matloob et al. [22]	Sequence Analysis, Rule Engine & Prediction-Based Engine	Reached an average accuracy of 85% and analyzed complex patient sequences across more than 62 specialties	Future research should test the framework on much larger datasets to better evaluate the model’s competence

mapping is formally outlined in Algorithm 1.

A. DATA SOURCES AND STUDY DESIGN

We used the Medicare dataset from Medical Sphere, which comes with 4 training and 4 test files. Both of them have provider, inpatient, and outpatient claims, as well as beneficiary tables. The training provider table includes a *Potential-Fraud* flag. Inpatient and outpatient tables contain rows with claim IDs, beneficiary IDs, provider IDs, dates, reimbursed amounts, deductibles, and diagnosis/procedure codes. The inpatient table also has admission and discharge dates and diagnosis-related fields. The beneficiary table contains basic patient information, demographics, and the patient’s chronic health condition. We computed mean, median, and z-scores on payments and length of stay. We also did the same for total claims, reimbursement amounts, and unique beneficiaries over 30/90/365 days to classify fraudulent providers. We utilized claims data for feature construction and performed unsupervised clustering on fraud-only claims. For supervised learning, we trained our models on 70% of the data and held out 30% for evaluation. We then applied the trained models to unlabeled test data to examine risk scores and labels for characterizing the provider clusters.




FIGURE 1. Experimental design diagram of the proposed methodology. The pipeline shows the input of Medicare data, the SQL-based feature engineering process, and the evaluation using supervised machine learning models and unsupervised HDBSCAN clustering.

Algorithm 1 Hybrid Healthcare Fraud Detection and Characterization Pipeline**Require:** Raw Medicare Data $\mathcal{D}_{raw} = \{\text{Claims}_{IP}, \text{Claims}_{OP}, \text{Beneficiaries}, \text{Providers}\}$

- Ground truth label set Y where $y_i \in \{0, 1\}$ (1 = Fraudulent Provider)

Ensure: Trained classification model M_{best} and optimal threshold τ_{opt}

- Provider fraud risk scores S_{risk}
- Identified claim clusters $C_{claims}^{IP}, C_{claims}^{OP}$ and provider clusters $C_{providers}$
- Association rules R_{assoc} for fraud typologies

```

1: Phase 1: SQL-Based Feature Engineering & Data Curating
2: Clean and standardize  $\mathcal{D}_{raw}$  into unified fact table  $F$ 
3: for each provider  $i$  in  $F$  do
4:   Compute rolling aggregates (30, 90, 365 days) for claim counts, patient volumes, and reimbursements
5: end for
6: for each claim  $c$  in  $F$  do
7:   Compute peer-normalized  $z$ -scores  $(\mu, \sigma)$ :
8:   if IP then
9:     Group by DRG  $\rightarrow Z_{reimb} = (X_c - \mu_{DRG}) / \sigma_{DRG}$ 
10:  else if OP then
11:    Group by Diagnosis Prefix  $\rightarrow Z_{reimb} = (X_c - \mu_{Prefix}) / \sigma_{Prefix}$ 
12:  end if
13:  Apply binary rule-based indicators  $I$ :
14:   $I_{duplicate} \leftarrow 1$  if exact match OR (date  $\pm 1$  day AND amount  $\pm 1\%$ )
15:   $I_{upcode} \leftarrow 1$  if ( $Z_{reimb} > 2.5$ ) AND ( $Z_{LOS} < -0.5$  OR  $LOS < 2$  days)
16:   $I_{overcharge} \leftarrow 1$  if ( $Z_{reimb} > 3.0$  OR  $X_{reimb} > Q3 + 1.5 \times IQR$ )
17: end for
18: Aggregate claim-level features and  $I$  flags to generate provider-level feature matrix  $X_{prov}$ 
19: Phase 2: Supervised Provider Fraud Classification
20: Join  $X_{prov}$  with labels  $Y$ 
21: Partition data:  $X_{train}, X_{test} \leftarrow \text{StratifiedSplit}(X_{prov}, 70/30)$ 
22: Apply Standard Normalization:  $X_{scaled} \leftarrow (X - \mu_{train}) / \sigma_{train}$ 
23: Apply class-weighting to mitigate label imbalance
24: Train classifiers  $\in \{\text{Logistic Regression (Elastic-Net), Random Forest, Extra Trees, LightGBM, XGBoost, CatBoost}\}$ 
25: Evaluate models using 5-fold Cross Validation optimized for AUPRC
26: Select  $M_{best} \leftarrow \text{Logistic Regression}$ 
27: Define  $\tau_{opt} \leftarrow \text{Threshold maximizing F1-Score on PR Curve}$ 
28: Generate probability scores:  $S_{risk} \leftarrow M_{best}.\text{predict\_proba}(X_{scaled})$ 
29: Phase 3: Unsupervised Fraud Typology Mapping
30: Claim-Level Analysis:
31: Filter  $X_{claims}$  where  $Y = 1$  (Fraud-only cases)
32: Apply HDBSCAN distinctly to Inpatient and Outpatient subsets  $\rightarrow C_{claims}^{IP}, C_{claims}^{OP}$ 
33: Run Apriori algorithm within clusters to extract diagnosis-procedure association rules  $R_{assoc}$ 
34: Provider-Level Analysis:
35: Filter  $X_{prov}$  where  $S_{risk} \geq \tau_{high\_risk}$  OR  $Y = 1$ 
36: Apply HDBSCAN ( $\text{min\_cluster\_size}=10$ ) to isolate high-risk provider groups  $C_{providers}$ 
37: Project  $C_{providers}$  into 2D space using Principal Component Analysis (PCA)
38: return  $M_{best}, \tau_{opt}, S_{risk}, C_{claims}^{IP}, C_{claims}^{OP}, C_{providers}, R_{assoc}$ 

```

B. DATA CLEANING AND SQL FEATURE STORE

All data processing was implemented in PostgreSQL using a layered design. We loaded raw CSV files into a staging schema with all fields as text and basic checks on structure and column counts. Cleaned data moved to a curated schema, where placeholder strings for missing values were converted to NULL and columns cast to appropriate types. We went through data

processing steps to verify unique provider and beneficiary IDs. We ensured that all features had non-null values across the inpatient and outpatient tables. We also looked into valid foreign keys and non-negative payments, length of stay, and other numeric features. After handling those checks, we put together a mart schema with a fact table and feature tables focused on claims and providers. These tables contain details

on length of stay, code counts for diagnoses and procedures, duplicates, possible overcharges, peer-normalized z-scores, and 30-, 90-, and 365-day aggregates.

C. FEATURE ENGINEERING AND RULE-BASED INDICATORS

We implemented simple rules on both claim and provider features to capture duplicate billing, upcoding, overcharging, and other fraud patterns. We worked on the reimbursed amount, deductible, length of stay, and counts of diagnoses and procedures, depending on the type of patient. We also detected exact or near duplicates by comparing the same patient and provider within a short time span. We then calculated peer-normalized z-scores and IQR thresholds for reimbursement payments and grouped patient records. To construct peer-normalized features, claims were grouped by their respective clinical categories (Diagnosis Related Group for inpatient claims and Primary Diagnosis Prefix for outpatient claims) across 30-, 90-, and 365-day windows. The standardized deviation Z for the reimbursement amount of a given claim c was mathematically formulated as:

$$Z_c = \frac{X_c - \mu_p}{\sigma_p} \quad (1)$$

where X_c represents the claim's reimbursement amount, while μ_p and σ_p represent the mean and standard deviation of all claims within the same peer group p . X_c values were winsorized at the 99th percentile before calculation to prevent extreme outliers from skewing the peer mean [17]. We then observed duplicates, upcoding, and overcharging to flag suspicious claims and used these inputs to train the models in the next steps.

D. EXPLORATORY ANALYSIS AND FEATURE SELECTION

After the initial setup of the SQL feature store was completed, we proceeded with the exploratory analysis and feature selection in Python. Using SQLAlchemy and pandas, we loaded the claim- and provider-level feature tables from the database to inspect the basic distribution of the data. To evaluate skewness and long tails in reimbursed amounts, deductibles, and inpatient length of stay, we plotted histograms and boxplots, which confirmed the expected tail behavior. We used IQR- and median-based z-scores to flag extreme outlier claims, then compared them with overcharge indicators and inpatient claims. A claim's reimbursement X_c was flagged as a hard overcharge outlier if it satisfied the following condition:

$$X_c > Q_3 + 1.5 \times (Q_3 - Q_1) \quad (2)$$

where Q_1 and Q_3 denote the 25th and 75th percentiles of the reimbursement amounts within the claim's respective peer group. We used pairwise correlation coefficients and inflation factors to remove highly collinear features, as those are redundant in the feature space. We were interested in the shared ground of recent reimbursement and claim volumes.

Therefore, the analysis enabled us to work on a meaningful set of provider features.

E. UNSUPERVISED FRAUD-ONLY CLAIM CLUSTERING

We clustered fraud-only claims to understand the diversity in fraudulent behaviors. We first distinguished fraudulent providers using the training labels and grouped all associated claims by inpatient and outpatient types. We stored those fraudulent records along with the extracted features in separate tables for inpatient and outpatient claims. We identified missing values and replaced them with the median across all records. We scaled features to ensure their contributions to the output were equal, which also sped up the optimization. It also prevented bias toward the larger numeric features. We then applied the HDBSCAN clustering algorithm distinctly to inpatient and outpatient claims [23]-[27]. Our goal was to find clusters of different sizes and densities, and claims that fell under the minimum number of clusters would be considered as noise [25]-[27]. We made sure the minimum size of clusters and samples were tuned for better accuracy. We also visualized the structure via PCA projections using different colors on the clusters. After that, we ran association rule mining within those clusters to characterize unique types of fraudulent claims.

To extract meaningful diagnosis-procedure linkages, the Apriori algorithm [11], [28]-[30] evaluates rules of the form $A \rightarrow B$. The strength and reliability of these rules were measured mathematically using Support, Confidence, and Lift, defined as:

$$\text{Support}(A \rightarrow B) = P(A \cap B) \quad (3)$$

$$\text{Confidence}(A \rightarrow B) = \frac{P(A \cap B)}{P(A)} \quad (4)$$

$$\text{Lift}(A \rightarrow B) = \frac{P(A \cap B)}{P(A) \times P(B)} \quad (5)$$

Rules were filtered based on a minimum Support of 0.02, a Confidence of 0.60, and a Lift threshold ($\text{Lift} > 1.2$) to statistically ensure that the identified fraudulent linkages happened more frequently than would be expected under random independence.

F. SUPERVISED PROVIDER FRAUD CLASSIFICATION

We applied supervised learning to classify providers as fraudulent or non-fraudulent based on the collected claim features. We trained the models only on labeled training providers, as we know these are absolutely fraudulent cases. We split the fraud cases from the training dataset into a 70/30 holdout set, using 70% for training and 30% for evaluation. This strategy preserved the proportion of fraudulent to non-fraudulent claims. We compared the performance using multiple classifiers on the same feature set. We used a linear model (logistic regression with elastic-net regularization) and tree-based models (Random Forests, Extra Trees, LightGBM,

CatBoost, XGBoost) for the benchmark evaluation on accuracy, recall, F1-score, AUPRC, and ROC-AUC [16]. Based on those models, we used class weighting or imbalance-aware parameters to address the scarcity of fraudulent providers. Our optimal-performing model was Regularized Logistic Regression, and it is important to define its objective function. To handle the severe class imbalance problem, we introduced a heuristic class weight (ω_i) into the cross-entropy loss function [16], [31]-[34]. Furthermore, we applied an Elastic-Net penalty to prevent overfitting across the high-dimensional feature space. The model optimized the coefficients (β) by minimizing the following loss function $L(\beta)$:

$$L(\beta) = -\frac{1}{N} \sum_{i=1}^N \omega_i [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)] + \lambda \left(\alpha \|\beta\|_1 + \frac{1 - \alpha}{2} \|\beta\|_2^2 \right) \quad (6)$$

where N is the number of providers, y_i is the true label, $\hat{y}_i$ is the predicted probability λ controls the overall regularization strength, and α determines the mix between L_1 (Lasso) and L_2 (Ridge) penalties. We tuned the hyperparameters using randomized search with 5-fold cross-validation and then optimized the AUPRC curve. We did not use the fixed 0.5 cutoff threshold, as it won't always yield the best result in this scenario. We chose the probability threshold that gave us the maximum F1-score on the AUPRC curve. After completing the performance tests, we found logistic regression to be the most performant model, and these performance gains have mostly been achieved by the engineered features rather than the model complexity.

G. PROVIDER FRAUD-TYPE CLUSTERING

To understand the diversity in fraud patterns, we went beyond classifying a provider as fraudulent or not; we sought to characterize fraud activities within the group of high-risk providers [29]. We characterized providers already labeled as fraudulent and test providers whose fraud scores have already passed the high threshold. We pulled out the same standardized features used in supervised models. These features were reimbursement amounts, claim volumes over an aggregated timeframe, the count of distinct patients, peer-normalized inpatient and outpatient z-scores, and rule scores for duplicate, upcoding, and overcharging claims. We then applied the HDBSCAN algorithm with a minimum cluster size of 10 and a small minimum sample value. This formed clusters with varying sizes and densities, and a few providers were scattered across the plot and were considered noise. We visualized the structure using principal components and colored points by cluster label, with noise in grey. For each cluster, we produced a labeled summary against high-scoring mix and interpreted those fraud types as high-revenue inpatient overchargers, high-volume outpatient overbillers, and others.

IV. RESULTS




FIGURE 2. Histogram and density plots of reimbursed amounts for inpatient and outpatient claims, showing a heavy right skew and a small number of very high-cost claims.

A. DATA CHARACTERISTICS AND EXPLORATORY ANALYSIS

The training data includes 40,474 inpatient claims, 517,737 outpatient claims, 138,556 distinct patients, and 5410 labeled providers, along with ground-level information on fraud or non-fraud. We found reimbursed amounts highly skewed, as expected in Medicare billing. In most cases, the claims had moderate payments, and a small fraction of claims were very costly, which produced a low-payment bulk and a long right tail (Figure 2). There were also variations in the inpatient and outpatient claims. Inpatient claims had a higher median reimbursed amount and a wider spread, confirming that inpatient services are more expensive and variable. In contrast, outpatient services show the claims count by length of stay, diagnosis vs. procedure counts, and reimbursement vs. coding intensity (Figure 3). We flagged high-cost claims using interquartile range and median absolute deviation rules. We kept a small set of potential outliers as possible fraud and unusual cases. After preprocessing the data, we found no missing values in the key features. Correlations and variance inflation factors (Figure 4) produced a compact feature set for unsupervised and supervised analyses.

B. UNSUPERVISED CLAIM-LEVEL PATTERNS AMONG FRAUDULENT CLAIMS

To understand the diversity of fraudulent behavior, we focused only on fraud-only cases. We identified 22,640 fraudulent inpatient and 8,408 fraudulent outpatient claims out of 558,211 claims. After standardizing the reimbursement amounts, code counts, rule flags, and peer-normalized scores, we implemented HDBSCAN on inpatient and outpatient records. HDBSCAN produced several clusters, and claims that did not belong to any cluster were considered noise. Figure 5 showed inpatient fraudulent clusters of high payments and long stays, as we obtained high reimbursement/length-of-stay z-scores. Figure 6 showed the visualizations of high-payment groups of




FIGURE 3. Boxplots of reimbursed amounts and inpatient length of stay, highlighting distributional differences between claim types.




FIGURE 4. Clustered Pearson correlation heatmap of provider-level features, used to identify highly collinear variables and guide feature selection.

fraudulent outpatient claims. Most of the outpatient clusters had low- to moderate-paid claims with very few higher-paid claims. Association rule mining within clusters (Table 2) showed small diagnosis-code sets in outpatient clusters and diagnosis-procedure pairs in high-payment inpatient clusters.

TABLE 2. Association rules for diagnosis and procedure code patterns within selected inpatient and outpatient HDBSCAN claim clusters. Results include antecedent code sets, consequent codes, support, confidence, and lift values for fraud-only clusters.

Claim Type	Cluster	Antecedents Str	Consequents Str	Support	Confidence	Lift
OP	1	280+285	588	0.039	0.766	10.127
OP	1	285+588	280	0.039	0.688	9.462
IP	18	995	38	0.031	0.645	8.896
OP	1	V56	588	0.024	0.637	8.425
OP	1	588	280	0.046	0.612	8.412
OP	1	280	588	0.046	0.636	8.412
IP	14	8154	715	0.031	0.863	7.415
IP	13	7935	820	0.051	0.865	7.398




FIGURE 5. Two-dimensional PCA projections of standardized fraud-only claims, showing HDBSCAN clusters of fraudulent inpatient claims. Points are colored by cluster labels, with noise points shown in grey.




FIGURE 6. Two-dimensional PCA projections of standardized fraud-only claims, showing HDBSCAN clusters of fraudulent outpatient claims. Points are colored by cluster labels, with noise points shown in grey.

C. PROVIDER-LEVEL FRAUD CLASSIFICATION PERFORMANCE

We evaluated supervised models for classifying fraudulent providers across all cases. On the 30% hold-out evaluation trainset, all models achieved ROC-AUC around 0.93-0.94, indicating that those models were good at telling apart fraudulent and non-fraudulent providers. All the performance metrics are reported in Table 3. Logistic regression with elastic-net regularization resulted in the best AUPRC of around 0.73 and ROC-AUC of around 0.94. After threshold tuning, we got a precision of around 0.59, a recall of around 0.74, and an F1-score of around 0.66. Tree-based models (Random Forests, LightGBM, XGBoost, Extra Trees, CatBoost) [16] achieved similar ROC-AUC but weaker AUPRC and could not surpass the logistic regression model (Figure 7). Fraud risk scores from the linear model placed most non-fraudulent providers near 0 and fraudulent providers at higher values (Figure 8). Big reimbursements, claim values, z-scores in the peer standard feature sets, and rule indicators of duplicates and overcharging contributed most to these high scores. We then applied those scores to the test providers to

define the high-risk providers for clustering in Section IV-D.

TABLE 3. Classification performance of supervised models on the provider hold-out set. Results are reported for logistic regression (elastic net), Random Forests, ExtraTrees, LightGBM, XGBoost, and CatBoost in terms of accuracy, precision, recall, F1 score, AUPRC, and ROC-AUC.

Model	Accuracy	Precision	Recall	F1-Score	AUPRC	ROC-AUC
LogReg	0.93	0.59	0.74	0.66	0.73	0.94
ExtraTrees	0.91	0.53	0.74	0.61	0.70	0.93
CatBoost	0.92	0.57	0.64	0.60	0.69	0.93
RandomForest	0.93	0.62	0.61	0.61	0.69	0.93
LightGBM	0.91	0.51	0.78	0.61	0.68	0.93
XGBoost	0.91	0.51	0.78	0.61	0.68	0.93




FIGURE 7. Bar charts of classification performance of supervised models tested for benchmark comparison.



FIGURE 8. Histogram shows the distribution of logistic regression fraud risk scores for providers. Scores for non-fraudulent providers cluster near zero, while scores for fraudulent providers shift toward higher values.

D. PROVIDER FRAUD-TYPE CLUSTERS

Among the high-risk class of 570 providers, HDBSCAN identified 4 clusters and labeled 334 providers as noise, meaning that many high-risk providers were not assigned to any cluster. Figure 9 shows compact regions for clusters with mixed noise. The largest cluster (Cluster 3) comprises 180 provider points, characterized by high z-scores for reimbursements and length of stay, but moderate counts of duplicates and overcharges. Cluster 1 (36 providers) comprises high-revenue providers with high inpatient z-scores, potentially indicating upcoding. Cluster 0 and Cluster 2 (10 providers each) account for the high count of outpatient billers with duplicates and overcharge flags. Table 4 shows the results of the fraudulent clusters.




FIGURE 9. Two-dimensional PCA projection of standardized provider-level features for high-risk providers. Points are colored by HDBSCAN cluster label, with noise providers shown in grey.

TABLE 4. Provider fraud-type clusters identified by HDBSCAN, showing the number of providers, labeled fraud cases in training, high-scored test providers, and median annual reimbursement and claim volumes.

cluster	n_providers	n_train_fraud	median_reimb_365d	median_claims_365d
0	10	10	252570	309
1	36	36	490880	140
2	10	10	390030	1336
3	180	180	210265	108

E. COMPARISON WITH THE STATE-OF-THE-ART METHODS

To validate the effectiveness of our proposed hybrid methodology, we compared our optimal performance against recent state-of-the-art (SOTA) studies focusing on healthcare fraud detection. While our literature review identifies a range of methodologies, this quantitative comparison strictly isolates benchmarks that yield supervised predictive metrics (e.g., ROC-AUC, AUPRC) for direct mathematical evaluation. Direct comparisons are challenging due to variations in datasets and class imbalance, yet we evaluated our metrics against benchmark studies that used similar Medicare claims data [16]. The comparative analysis prioritizes Area Under the Receiver Operating Characteristic Curve (ROC-AUC) and the Area Under the Precision-Recall Curve (AUPRC) over single-point metrics such as Precision, Recall, and the F1-score. F1-scores provide insight into model performance at an arbitrary classification threshold (typically 0.5). They do not capture an algorithm's predictive ability across diverse environments. In contrast, ROC-AUC and AUPRC are threshold-agnostic [17]. They are aggregate metrics that measure the total area under the curve across all possible thresholds from 0.0 to 1.0. As we are dealing with the severe class imbalance inherent in Medicare claims data, AUPRC is regarded as the critical evaluation metric in the study, since it is evaluated without accounting for True Negatives [17]. Also, it provides a measure of the model's ability to accurately isolate the rare minority class (fraud) without skewing [17].

As shown in Table 5, the proposed SQL-based feature engineering with a class-weighted Regularized Logistic Regression achieves a highly competitive ROC-AUC of 0.94

TABLE 5. Performance Comparison with State-of-the-Art Fraud Detection Models

Study	Methodology	Dataset	Best ROC-AUC	Best AUPRC
Nabrawi & Alanazi [13]	Tree-based models & ANN	Saudi Insurance Claims	90%	Not Reported
Hancock et al. [17]	Ensemble Selection & GBDTs	Medicare Part B & D	95.88%	0.8016
Kanksha et al. [19]	Isolation Forest	Medicare Claims	94.32%	Not Reported
Proposed Approach	SQL-Feature Store + Regularized LR	Medicare Claims	94%	0.73

and an AUPRC of 0.73. Our approach almost matches or outperforms the predictive capabilities of complex gradient boosting ensembles (GBDTs) proposed in recent literature. For example, our AUPRC of 0.73 outperforms the maximum AUPRC (0.71) reported by Hancock et al. [17] on Medicare Part B data, despite our use of a highly interpretable linear model. Furthermore, our results emphasize the necessity of Precision-Recall metrics over standard Accuracy. While studies like Kanksha et al. [19] report accuracies exceeding 98%, such metrics often fail to capture minority-class detection in severely imbalanced fraud domains.

V. CONCLUSION AND FUTURE WORK

This study shows that a comprehensive and hybrid pipeline provides an effective framework for identifying and characterizing complex Medicare fraud. Our results pertain to classical machine learning methods, specifically regularized logistic regression, which outperforms complex tree-based models. Our strategy follows contemporary research combining association rules, supervised models, and unsupervised clustering, rather than relying on a single model [1]-[4]. The regularized logistic regression model achieved an exceptional ROC-AUC of 0.94 and an AUPRC of 0.73, indicating that prediction of fraudulent cases is driven more by feature engineering than by algorithm complexity [1], [2], [13]. Furthermore, our analysis goes beyond binary classification by successfully isolating distinct fraud typologies. Deep learning offers additional performance gains from a healthcare perspective but requires greater computational power and resources [15]. This research offers a practical benchmark for real-world investigative systems by translating complex claims data into actionable clusters such as high-revenue inpatient overcharging or high-volume outpatient overbilling [12], [20]. Ultimately, integrating a structured feature store with both supervised and unsupervised learning enables healthcare organizations to develop explainable, high-performance tools to combat sophisticated financial abuse, thereby following best practices to restore integrity to public insurance programs [3], [4].

We have encountered a class imbalance problem, as most cases are labeled non-fraudulent, and our models are becoming biased toward normal cases. The complexity in feature engineering is profound. We tried to resolve the class imbalance problem using 50% of the fraudulent and 50% of the non-fraudulent cases from the train dataset. Also, we had computational limitations, and so we did not integrate the deep learning models, as it would have cost a lot of time and power. The dataset we worked on covered only one year

of patients' data, which might not be sufficient to obtain a holistic result. If we could work on the dataset in the expanded timeframe, the models could produce more mature output. We could not verify the flagged cases as fraudulent, as we need to verify that further with the Subject Matter Expert/Healthcare Professional. Currently, the approach is train-test split-based, and so it does not provide any real-time fraud scoring on patients' data.

We can further enhance those models' performance by effectively tuning hyperparameters or by focusing on SQL-based feature engineering to improve performance metrics. Most studies have confirmed that SQL feature engineering is a crucial step that plays a significant role in improving evaluation scores [1], [2], [13]. We can also collaborate with a healthcare professional to assess the model's output for flagging fraudulent cases by validating the association rule mining on those claims. Subject-matter experts can review the top 100 claims or providers from the fraud-type clusters and help us refine the rule definitions. They can also help us ensure which clusters belong to which actionable schemes and provide more reliable labels. We can definitely test on the multi-year dataset, as it will provide more insights for fraud detection and can be integrated with deep learning models, despite being computationally intensive. Then we can move on to building a real-time fraud-scoring system. We will embed the current SQL feature store and the model training into the streaming architecture. New claims will automatically be appended to the staging layer; all the aggregates will be updated near real time. Based on the training set, the model will assign fraud scores to new claims and providers as they arrive. We can also integrate live dashboards where we can see the claims flagged as fraudulent for suspicious behavior and non-fraudulent cases as well. Integrating all of the above would give rise to an innovative system for capturing atypical claims and providers in the healthcare industry, ensuring greater transparency.

ACKNOWLEDGMENT

The authors would like to express their sincere gratitude to Dr. Sarika Rajeev and Dr. Timothy Oladunni for their valuable guidance, insightful suggestions, which significantly improved the quality of this manuscript.

REFERENCES

- [1] A. du Preez, S. Bhattacharya, P. Beling, and E. Bowen, "Fraud detection in healthcare claims using machine learning: A systematic review," *Artif. Intell. Med.*, vol. 160, art. 103061, 2025, doi: 10.1016/j.artmed.2024.103061.
- [2] N. Kumaraswamy, M. K. Markey, T. Ekin, J. C. Barner, and K. Rascati, "Healthcare fraud data mining methods: A look back and look ahead," *Perspect. Health Inf. Manag.*, vol. 19, no. 1, p. 1i, 2022.

- [3] K. Razzaq and M. Shah, "Next-generation machine learning in healthcare fraud detection: Current trends, challenges, and future research directions," *Information*, vol. 16, no. 9, art. 730, 2025, doi: 10.3390/info16090730.
- [4] A. Y. B. R. Thaifur, M. A. Maidin, A. I. Sidin, and A. Razak, "How to detect healthcare fraud? A systematic review," *Gac. Sanit.*, vol. 35, suppl. 2, pp. 441–449, 2021, doi: 10.1016/j.gaceta.2021.07.022.
- [5] M. Bairy, B. Muniyal, and N. P. Shetty, "Enhancing healthcare data integrity: Fraud detection using unsupervised learning techniques," *Int. J. Comput. Appl.*, vol. 46, no. 11, pp. 1006–1019, 2024, doi: 10.1080/1206212X.2024.2408262.
- [6] M. Herland, T. M. Khoshgoftaar, and R. A. Bauder, "Big Data fraud detection using multiple Medicare data sources," *Journal of Big Data*, vol. 5, no. 1, p. 29, 2018.
- [7] R. A. Bauder and T. M. Khoshgoftaar, "Medicare fraud detection using machine learning methods," in *16th IEEE International Conference on Machine Learning and Applications (ICMLA)*, pp. 858–865, 2017.
- [8] K. Mon, "Medicare Fraud Detection: Data Analytics Approach," *NORMA@NCI Library*, 2024.
- [9] A. Abdallah, M. A. Maarof, and A. Zainal, "Fraud detection system: A survey," *Journal of Network and Computer Applications*, vol. 68, pp. 90–113, 2016.
- [10] F. Ahammed, B. A. Barakati, and O. Marques, "Improving Medicare Fraud Detection Accuracy in Deep Learning by Exploring Feature Selection and Data Sampling Techniques," *medRxiv*, Mar. 2026.
- [11] Z. Hamid, F. Khaliq, S. Mahmood, A. Daud, A. Bukhari, and B. Alshemaimri, "Healthcare insurance fraud detection using data mining," *BMC Med. Inform. Decis. Mak.*, vol. 24, art. 112, 2024, doi: 10.1186/s12911-024-02512-4.
- [12] M. C. Massi, F. Ieva, and E. Lettieri, "Data mining application to healthcare fraud detection: A two-step unsupervised clustering method for outlier detection with administrative databases," *BMC Med. Inform. Decis. Mak.*, vol. 20, art. 160, 2020, doi: 10.1186/s12911-020-01143-9.
- [13] E. Nabrawi and A. Alanazi, "Fraud detection in healthcare insurance claims using machine learning," *Risks*, vol. 11, no. 9, art. 160, 2023, doi: 10.3390/risks11090160.
- [14] Q. Liu and M. A. Vasarhelyi, "Healthcare fraud detection: A survey and a clustering model incorporating geo-location information," in *Proc. 29th World Continuous Auditing and Reporting Symp. (WCARS)*, Brisbane, Australia, Nov. 21–22, 2013.
- [15] P. A. Kumar and S. Sountharajan, "Insurance claims estimation and fraud detection with optimized deep learning techniques," *Sci. Rep.*, vol. 15, art. 27296, 2025, doi: 10.1038/s41598-025-12848-0.
- [16] R. Bounab, K. Zarour, G. Bouchra, and N. Khelifa, "Enhancing Medicare Fraud Detection through Machine Learning: Addressing Class Imbalance with SMOTE-ENN," *IEEE Access*, vol. 12, pp. 54399–54413, 2024, doi: 10.1109/ACCESS.2024.3385781.
- [17] J. T. Hancock, R. A. Bauder, H. Wang, and T. M. Khoshgoftaar, "Explainable machine learning models for Medicare fraud detection," *Journal of Big Data*, vol. 10, no. 1, art. 154, 2023.
- [18] H. De Meulemeester, F. De Smet, J. van Dorst, E. Derroite, and B. De Moor, "Explainable unsupervised anomaly detection for healthcare insurance data," *BMC Med. Inform. Decis. Mak.*, vol. 25, art. 14, 2025, doi: 10.1186/s12911-024-02823-6.
- [19] Kanksha, A. Bhaskar, S. Pande, R. Malik, and A. Khamparia, "An intelligent unsupervised technique for fraud detection in health care systems," *Intell. Decis. Technol.*, vol. 15, no. 1, pp. 127–139, Feb. 2021, doi: 10.3233/IDT-200052.
- [20] L. Settipalli and G. R. Gangadharan, "WMTDBC: An unsupervised multivariate analysis model for fraud detection in health insurance claims," *Expert Syst. Appl.*, vol. 215, art. 119259, 2023, doi: 10.1016/j.eswa.2022.119259.
- [21] G. Saldamli et al., "Health care insurance fraud detection using blockchain," in *Proc. 7th Int. Conf. Softw. Defined Syst. (SDS)*, 2020, pp. 145–152, doi: 10.1109/SDS49854.2020.9143900.
- [22] I. Matloob, S. A. Khan, and H. U. Rahman, "Sequence mining and prediction-based healthcare fraud detection methodology," *IEEE Access*, vol. 8, pp. 143256–143273, 2020, doi: 10.1109/ACCESS.2020.3013962.
- [23] L. McInnes, J. Healy, and S. Astels, "hdbscan: Hierarchical density based clustering," *Journal of Open Source Software*, vol. 2, no. 11, p. 205, 2017.
- [24] R. J. G. B. Campello, D. Moulavi, and J. Sander, "Density-based clustering based on hierarchical density estimates," in *Pacific-Asia Conference on Knowledge Discovery and Data Mining*, pp. 160–172, Springer, 2013.
- [25] R. Setiawan, B. Tjahjono, G. Firmansyah, and H. Akbar, "Fraud Detection in Credit Card Transactions Using HDBSCAN, UMAP and SMOTE Methods," *International Journal of Science, Technology & Management*, 2024.
- [26] M. M. Breunig, H.-P. Kriegel, R. T. Ng, and J. Sander, "LOF: identifying density-based local outliers," in *ACM SIGMOD Record*, vol. 29, no. 2, pp. 93–104, 2000.
- [27] C. C. Aggarwal, "Outlier analysis," in *Data mining*, pp. 237–263, Springer, 2015.
- [28] R. Agrawal and R. Srikant, "Fast algorithms for mining association rules," in *Proc. 20th int. conf. very large data bases, VLDB*, vol. 1215, pp. 487–499, 1994.
- [29] S. Zhou, J. He, H. Yang, D. Chen, and R. Zhang, "Big Data-Driven Abnormal Behavior Detection in Healthcare Based on Association Rules," *IEEE Access*, vol. 8, pp. 129002–129011, 2020, doi: 10.1109/ACCESS.2020.3009006.
- [30] J. Liu, "Graph Analysis for Detecting Fraud, Waste, and Abuse in Healthcare Data," *AAAI*, vol. 29, no. 2, pp. 3912–3919, Jan. 2015.
- [31] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: Synthetic minority over-sampling technique," *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.
- [32] J. M. Johnson and T. M. Khoshgoftaar, "Medicare fraud detection using neural networks," *Journal of Big Data*, vol. 6, no. 1, p. 63, 2019.
- [33] A. Fernandez, S. Garcia, F. Herrera, and N. V. Chawla, "SMOTE for learning from imbalanced data: progress and challenges, marking the 15-year anniversary," *Journal of Artificial Intelligence Research*, vol. 61, pp. 863–905, 2018.
- [34] T. Khoshgoftaar et al., "Impact of Class Imbalance on Unsupervised Label Generation for Medicare Fraud Detection," *NSF Public Access*, 2024.



ing, and Data Science.

ZAYED M. UDDIN received the B.S. with Summa Cum Laude from the Department of Mathematics & Computer Science, Coppin State University, Baltimore, USA, in 2026 and he will be pursuing the Ph.D. degree with the Department of Computer Science, University of North Carolina in Fall 2026. He has been awarded the best academic distinction award across the department. His research interests include Deep Learning (vision transformers; CNN + attention for object detection), Machine Learn-



MD. ABU BAKAR SIDDIK MISHKAT is currently pursuing B.S. in Computer Science with the Department of Mathematics & Computer Science at Coppin State University, Baltimore, USA. He has been awarded prestigious scholarships from BGE and Presidio for his academic merit. His research interests include Virtual Reality, Augmented Reality, Crypto and Blockchain technologies.



ULUGBEK ABDURAKHMONOV received the B.S. with Summa Cum Laude from the Department of Mathematics & Computer Science, Coppin State University, Baltimore, USA, in 2026 and he will be pursuing the Ph.D. degree with the Department of Computer Science, Montana State University in Fall 2026. He was a four-year NCAA Division I student-athlete on the tennis team. His research interests include Cancer Detection using Computer Vision, System Design, Algorithms, and Artificial

Intelligence.

• • •